
What Counts as a Mistake? Annotating Recitation Events in Quran Memorization Transcripts

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Checking Quran recitation from an ASR transcript requires distinguishing unre-
2 solved mistakes from repetitions, repairs, opening formulas and accepted spelling
3 differences. We report a completed human annotation of 100 production recording
4 cases, covering 348 scored units and 162 localized events across ten combined
5 labels. An executable evaluator scores labels and word positions together. A plain
6 diff reaches label-aware F1 0.531 and localization F1 0.839; adapted production
7 cleaner/alignment components reach 0.518 and 0.799. Their exact-span F1 scores
8 are both 0.505. Correcting the adapter’s word coordinates recovers all five an-
9 notated repetition events, showing why annotation interfaces must be checked
10 before interpreting baseline failures. The dataset and scorecard support a future
11 experiment in agent-assisted annotation and algorithm development.

12 1 The distinction the task turns on

13 Quran memorization is normally practised with a listener who corrects the reciter. An application
14 that stands in for that listener transcribes with commercial ASR and then compares the transcript
15 against the reference. Edit distance locates differences but does not determine their interpretation.
16 In our production data 74% of reviewed ayah units carry no event at all, and among those that do, a
17 substantial share are things a human listener would never mention: the reciter repeated a phrase after
18 a breath, said a word wrongly and immediately fixed it, or used a written form differing from the
19 Uthmani orthography without differing in what was recited. Flagging these as unresolved mistakes
20 can give the learner misleading feedback; we do not measure the relative user cost here.

21 We therefore annotate transcript-against-reference differences as typed, localized *events*, and we
22 judge the text rather than the speaker: a label records what differs between transcript and reference,
23 not whether the reciter or the ASR caused it. Unwritten vowel and tajweed errors lie outside a
24 transcript-only rubric and outside this work.

25 2 Taxonomy and annotation

26 **Combined labels.** An earlier pilot used two parallel fields, an event type and a be-
27 nign/mistake verdict, plus an uncertain escape hatch. Reviewing real recordings made
28 both untenable: the verdict was never independent of the type in practice, and uncertain
29 absorbed exactly the cases that needed a decision. The rubric now uses one com-
30 bined label per event, `substitution_mistake`, `omission_mistake`, `insertion_mistake`,
31 `substitution_corrected`, `omission_corrected`, `repetition_benign`, `letters_benign`,
32 `spelling_benign`, `basmala_benign` and `isti3adha_benign`, with `clean` for a reviewed unit

33 carrying no event. A unit reviewed and found clean is not the same as a unit not yet reviewed, and the
34 data distinguishes them.

35 **Four decisions that took adjudication.** *Omissions* include missing beginnings, interiors and endings,
36 with no exemption for the final ayah of a recording: stopping early still leaves reference material
37 absent from the transcript, and a reviewer cannot infer intent from recording position. *Repeats*
38 span the original phrase together with its extra occurrence against a single reference copy, aligning
39 their localization with corrected events; a wrong-then-correct attempt is a corrected event, while
40 a correct-then-incorrect restatement is a *substitution_mistake* spanning both attempts, with
41 the order recorded in the note. *Spelling* differences removed by an agreed context-independent
42 normalization receive no event, and only a residual form needing contextual acceptance receives
43 *spelling_benign*; the accepted contexts are enumerated rather than generalized, so the category
44 cannot quietly become blanket final-letter deletion. *Localization* uses verbatim words and half-open
45 spans into the original whitespace token arrays, with an empty hypothesis span anchoring an omission
46 and an empty reference span anchoring an insertion.

47 **Scale and composition.** All 100 selected recording cases are human-approved: 314 ayah units,
48 233 of them clean, 115 within-ayah events and 47 opening-formula annotations, over 274 ayahs
49 in 38 surahs. The event distribution is uneven by nature, from 62 *substitution_mistake* down
50 to one *insertion_mistake* and one *omission_corrected*, and any evaluation has to survive
51 that. Ten learner identities contribute, one of them 46 recordings. Selection was purposive, so these
52 counts describe the annotated set; they do not estimate error prevalence or imply generalization to
53 unseen reciters. Three further approved recordings are held as reserves outside the scored set. We
54 verified identifiers and transcripts against both source exports and excluded punctuation-normalized
55 duplicates and redundant clean subpassages. Gold answers and selected membership remain in an
56 owner-held private bundle; this draft does not release a gold sample.

57 3 An executable scorecard

58 Systems receive the reviewed ayah split, the ayah ids and the exact references, so detection and
59 segmentation are supplied rather than scored; that task is treated in a companion submission. Per unit
60 a system returns events of the form {label, hyp_span, ref_span}, with half-open spans into
61 the transcript and reference token arrays.

62 **Opening formulas are scored.** Each record’s isti’adhah and basmala are presented as their own unit,
63 with the formula as the transcript and an empty reference; a compound opening yields two events
64 over one token array. Deciding that a basmala preceding the ayah is benign rather than an insertion is
65 exactly the judgment under test, and excluding it would hide the failure mode that turns out to matter
66 most.

67 **Matching.** Predictions and gold events are paired one-to-one within a unit by maximum total
68 similarity, so a single broad prediction is credited with at most one gold event and the remainder count
69 as misses. A pair is eligible only when *both* spans overlap, $\min(\text{sim}_{hyp}, \text{sim}_{ref}) \geq 0.30$; averaging
70 the two, as an earlier version did, lets an exact reference span carry a match while the hypothesis
71 span points somewhere else entirely. Events are validated for known labels, integer endpoints, token
72 bounds and label-specific empty spans. Invalid predictions count as false positives. Empty spans
73 match only empty anchors, exactly or within one token; they never match selected words. Pairing
74 ignores labels so localization can be reported separately. Larger units with more than seven events on
75 either side use deterministic greedy one-to-one matching.

76 **Scores.** The primary measure is label-aware event F1: a paired prediction is a true positive only
77 when its label also matches, and a mislabeled pair counts once as a false positive and once as a
78 false negative. We report micro F1 as the headline, exact-span F1 requiring identical endpoints,
79 macro F1 across labels present in gold, and localization precision/recall/F1 with labels dropped. An
80 application-facing cost is reported separately in raw event counts, (2 missed + false flags) per 100
81 units, so that the stated exchange rate of one missed mistake for two unnecessary flags is the one
82 the formula implements. The 2:1 ratio is a stated preference, not an empirically established user
83 cost. Predicting nothing gives F1 0 and the gold annotation gives 1. Evaluator v2.1 passes synthetic
84 coordinate, label, anchor and adapter regression tests and private gold oracle checks. The tolerance of
85 0.30, one-token anchor slack and secondary cost remain provisional until disjoint practice review;
86 baseline results are conditional on these choices.

87 4 What existing systems do

System	micro F1	exact F1	macro F1	loc F1	review cost
Predict nothing	0.000	0.000	0.000	0.000	51.2
Plain diff	0.531	0.505	0.179	0.839	22.4
Adapted production components	0.518	0.505	0.284	0.799	26.7

Table 1: Evaluator v2.1: 348 units, 162 gold events. All baseline calls complete without crashes or invalid predictions.

88 The plain diff locates 128 of 162 gold events (recall 0.790), with localization F1 0.839. Its label-aware
 89 recall is 0.500 and F1 is 0.531. These quantities differ: localization F1 is not the percentage of gold
 90 events found. It reaches per-label F1 0.94 for omissions and 0.85 for substitutions, and zero for the
 91 other eight labels. It emits 46 insertion predictions, 34 of them on opening-text units, whereas gold
 92 has one insertion event. Compound openings contain two gold formulas but a diff can return one
 93 insertion spanning both.

94 The second baseline adapts existing production cleaner and alignment components to the new event
 95 schema; the deployed application does not natively expose it. The adapter retains original token
 96 identities through deletion, uses alignment indices and reference offsets, and maps cleaner notes to
 97 reference phrases and both attempts. Its explicit heuristic interprets a stumble as a repair only when
 98 the following equally long surviving span matches the reference; unlocalizable removed fragments
 99 remain extra text. Production’s partial-start scoring policy is retained. These choices are part of the
 100 baseline, not additional gold decisions.

101 This adapter predicts five repetitions and recovers all five (F1 1.0); six corrected-substitution predic-
 102 tions recover one of four gold events (F1 0.20). It has nonzero F1 on four labels and zero on six,
 103 and never predicts either opening-formula label, letters, contextual spelling or omission corrections.
 104 The earlier adapter lost valid repetition credit through incorrect coordinates. The remaining label
 105 gaps motivate explicit event modeling, but these results do not establish that cleaning itself causes the
 106 errors or measure every feature of the deployed application. Each baseline retains its own comparison
 107 normalization, which is an algorithm choice rather than a change to the gold rubric.

108 5 Limitations and what comes next

109 The annotated set is purposive and small, one learner contributes 46 of the 100 recordings, and the
 110 label distribution is heavily skewed, so macro F1 here is a coverage check rather than a stable estimate.
 111 More consequentially, all 100 selected transcripts already appear as inputs in a public release of
 112 the same recordings, and twelve appear in an earlier machine-labeled pilot. The adjudications and
 113 the selection membership are private, but the inputs are not, and runtime isolation cannot establish
 114 absence of prior exposure. Any claim about unseen inputs requires a separately collected corpus;
 115 a read-only production check has since identified 63 further recordings with 62 stored transcripts,
 116 forming 57 distinct text groups absent from the old export, as candidates, subject to reuse permission
 117 and de-identification.

118 The annotation and the scorecard exist to support a measured experiment we have not yet run: supply
 119 an agent with an unlabeled development pool, the reference and the label definitions, let it produce
 120 its own annotations and an algorithm, and score only the resulting algorithm against the private set.
 121 Agreement with self-generated labels would not be independent annotation accuracy, and the score
 122 assesses the combined process rather than any single annotation’s contribution. Budgets, comparison
 123 conditions, matching tolerances and permitted feedback remain to be frozen. Outside gold100 and
 124 its three reserves, 217 transcripts remain; a preliminary text screen leaves 166 distinct candidate
 125 groups, 65 requiring overlap review. Assign groups to separate development or future-test roles
 126 before sharing. The developing agent sees neither gold inputs nor answers, selection membership,
 127 reviewer files or annotation-session memory. Only the frozen algorithm receives gold input in
 128 an owner-controlled evaluation environment; removing answers alone does not make gold inputs
 129 suitable for development. Those private materials stay outside development runtimes, and shared
 130 documentation carries aggregate findings and synthetic examples only.